RITIK SONTAKKE

Agentic AI Engineer

CONTACT

+91 8668851662
ritiksontakke.work@gmail.com
GitHub
LinkedIn
X

IT SKILLS

Agentic AI & GenAI: Multi-Agent Orchestration, LangGraph, LangChain, Model Context Protocol (MCP), RAG, LLM Evaluation (Ragas), Tool Calling

Back-end & Data: Python, FastAPI, Spring Boot, SQL, Vector DBs (pgvector, Chroma), PostgreSQL,Redis

Cloud & DevOps: AWS (ECS, Lambda), Docker, Kubernetes (k8s), Terraform, GitHub Actions, CI/CD

CORE COMPETENCIES

- RAG Architecture
- LLM Evaluation & Optimization
- Back-end Engineering (Python)
- Cloud-Native Development (AWS,Docker, k8s)
- Vector Database Engineering (pgvector, Chroma)
- Microservices & API Design
- Secure Multi-Tenant Architecture
- CI/CD & DevOps Automation
- Cross-Functional Collaboration

EDUCATION

AI Engineering Specialization – *MisogiAI by Masai, Bengaluru*
Jun 2025 – Sept 2025

- Focus: AI/ML, Autonomous Agents, LangGraph, RAG
- Outcome: Designed and implemented **complex, multi-step tool-calling workflows**

Bachelor of Commerce – *Sant Gadge Baba Amravati University, Maharashtra*
Apr 2021 – Mar 2023

PERSONAL DETAILS

Address	: Amravati, Maharashtra
Date of Birth	: 17th July 2002
Languages Known	: English, Marathi & Hindi

ABOUT ME

To leverage expertise in backend and **Agentic AI engineering, including autonomous** agent workflows, **LangChain | LangGraph** orchestration, and Retrieval-Augmented Generation (RAG) systems, to design and deploy secure, scalable, cloud-native AI solutions that enable intelligent automation and reliable, data-driven decision-making in production environments.

PROFILE SUMMARY

- **Backend & Agentic AI Engineer with 1+ year of professional experience designing and deploying production-grade AI-powered backend systems.**
- Strong expertise in building Retrieval-Augmented Generation (RAG) pipelines and multi-step agent workflows using LangChain and LangGraph.
- Proven ability to design modular backend architectures with clean routing, dependency injection, and scalable service patterns using FastAPI and PostgreSQL.
- Hands-on experience implementing secure, multi-tenant cloud-native systems on AWS with Docker, Terraform, and CI/CD automation.
- Experienced in integrating LLM tool-calling frameworks, prompt templates, and context-aware AI pipelines for enterprise applications.
- Built document ingestion pipelines for PDF, DOC, and HTML sources, including text extraction, chunking strategies, embedding generation, and metadata storage.
- Comfortable working across the full backend lifecycle—from API development and database design to AI pipeline integration and deployment support
- Strong focus on production readiness through structured logging, centralized error handling, API versioning, and system observability.
- Hands-on experience in deploying **scalable cloud-native infrastructure** using Terraform, Docker, AWS ECS/Lambda, and CI/CD pipelines—drastically improving deployment efficiency and operational resilience.

WORK EXPERIENCE

Backend | AI Engineer | Edifition | Feb 2024 – Apr 2025 | Remote

- Designed and deployed scalable REST APIs using FastAPI, Python, and PostgreSQL to support AI-powered backend services.
- Built modular backend architecture with clean routing, dependency injection, middleware, and versioned API design.
- Implemented multi-step conversational workflows and agent supervision using LangGraph.
- Integrated LangChain for LLM tool-calling, prompt templates, and context-aware AI pipelines.
- Developed enterprise-grade RAG pipelines using vector databases to deliver accurate, document-grounded AI responses.
- Built document ingestion systems for PDF, DOCX, and HTML sources, including text extraction, chunking strategies, embedding generation, and metadata storage.
- Implemented structured logging, centralized error handling, and production-grade exception management.
- Designed schema-per-tenant PostgreSQL architecture to ensure strict data isolation and secure multi-client deployments.
- Integrated third-party APIs with robust retry logic, timeout handling, and failure recovery mechanisms.
- Supported cloud deployments using Docker and AWS, improving system scalability an operational reliability.

PROJECT

Enterprise RAG Orchestration Platform

- Architected a scalable **multi-format RAG pipeline** enabling enterprise-grade conversational intelligence with robust ingestion, indexing, and retrieval workflow.
- Implemented schema-per-tenant architecture in PostgreSQL, ensuring strict data isolation, security, and scalability across multiple client organizations.
- Integrated LangChain with OpenAI and Google Generative AI APIs to deliver accurate, context-aware responses grounded exclusively in tenant-specific documents.
- Architected and deployed the platform on AWS using Docker, ensuring scalability, fault tolerance, and cloud-native best practices.
- Designed a secure **Multi-Tenant Architecture** on AWS using schema-per-tenant PostgreSQL and RBAC, ensuring complete context isolation and preventing cross-tenant information leakage.
- Structured the backend with decoupled RAG services, enabling independent scaling, maintenance, and future extensibility.
- Developed an end-to-end RAG pipeline covering document upload (PDF, DOCX, TXT), asynchronous ingestion, chunking, vectorization using pgvector, and context-grounded AI response generation.
- Built a **high-performance agent orchestration system using LangGraph and LangChain**, integrated with pgvector-based hybrid vector search to reduce retrieval latency and enable real-time, context-aware agent reasoning.